

# GraphRAG and Knowledge Integration

Week 8: From Vector Search to Knowledge Graphs

PhD Course in Agentic Artificial Intelligence

12-Week Research-Level Course

## Bloom's Taxonomy Levels Covered

- **Remember:** Define knowledge graph (entity-relationship structure), entities, relations, communities (clusters)
- **Understand:** Explain how GraphRAG enhances retrieval with structure
- **Apply:** Build a knowledge graph from unstructured text using LLMs
- **Analyze:** Compare vector-only vs graph-enhanced retrieval strategies
- **Evaluate:** Assess when GraphRAG provides value over standard RAG
- **Create:** Design a hybrid retrieval system combining vectors and graphs

---

By end of lecture, you will understand structured knowledge integration in agents.

# Limitations of Vector-Only RAG

## What Vector Search Does Well

- Semantic similarity matching (“find documents about X”)
- Fast approximate nearest neighbor search
- Works with any text without preprocessing

## Where Vector Search Struggles

- **Multi-hop reasoning:** “Who founded the company that acquired X?”
- **Global queries:** “What are the main themes across all documents?”
- **Relationship traversal:** “How are entities A and B connected?”
- **Aggregation:** “List all products mentioned with their features”

## The Insight

- Structure enables reasoning that similarity alone cannot support

---

GraphRAG adds explicit structure to enable complex reasoning patterns.

# What is a Knowledge Graph?

## Definition

- A knowledge graph represents information as **entities** (nodes) connected by **relations** (edges)
- Triples: (subject, predicate, object) – e.g., (Apple, founded\_by, Steve Jobs)

## Key Components

- **Entities:** People, places, concepts, events (named objects)
- **Relations:** Connections between entities (typed edges)
- **Attributes:** Properties of entities (metadata)
- **Schema:** Optional type hierarchy and constraints

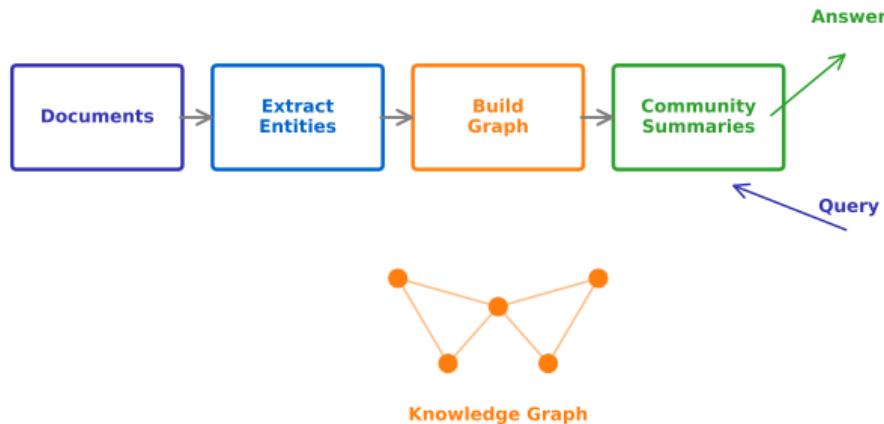
## Examples

- Wikidata (90M+ entities), Google Knowledge Graph, enterprise KGs

---

Knowledge graphs make implicit relationships explicit and queryable.

## GraphRAG: Knowledge Graph + RAG



GraphRAG builds structure from documents before retrieval.

# Entity and Relationship Extraction

## LLM-Based Extraction (Microsoft GraphRAG)

- Use LLM to identify entities and relationships from text chunks
- Prompt: “Extract all entities and their relationships from this text”
- Output: Structured triples (entity1, relation, entity2)

## Entity Types Commonly Extracted

- People, organizations, locations, events, concepts
- Domain-specific: products, chemicals, diseases, etc.

## Challenges

- Entity resolution (“Apple Inc.” = “Apple” = “the company”)
- Relation normalization (“works for” = “employed by”)
- Extraction quality depends heavily on LLM capability

---

LLM-based NER (Named Entity Recognition) enables extraction without training.

## Extraction Precision vs. Recall

- High precision: Miss valid entities/relations
- High recall: Include noise, false positives
- LLM extraction typically favors recall

## Schema Design Choices

- **Open schema:** Flexible but harder to query
- **Fixed schema:** Constrained but consistent
- **Hybrid:** Core schema + open extensions


## Cost Breakdown (per 1M tokens)

- Entity extraction: \$1-3, Relation extraction: \$2-5
- Entity resolution: \$0.5-1, Community summarization: \$3-8
- Total GraphRAG indexing: ~\$10/1M tokens vs. \$0.10 for basic RAG

---

GraphRAG indexing cost is 100x basic RAG – ensure query patterns justify it.

# Entity Extraction Pipeline



### Extraction Stats:

Entities: ~50/chunk

Relations: ~30/chunk

---

LLMs extract entities and relations to build the knowledge graph.

## The Global Query Problem

- Local queries: “What did X say about Y?” – direct retrieval works
- Global queries: “What are the main themes?” – needs aggregation

## Hierarchical Summarization via Communities

- **Leiden algorithm** (graph clustering method): Cluster densely connected entities
- Generate summaries at each community level
- Build hierarchy: documents → entities → communities → global

## Benefits

- Pre-computed summaries enable fast global queries
- Multiple granularity levels for different query types
- Captures both local detail and global themes

---

Community summaries enable answering “what is this corpus about?” efficiently.

## Hierarchical Community Detection



### Community Summaries (LLM-generated)

C1: "AI models..."

C2: "Training data..."

C3: "Applications..."

Leiden algorithm clusters entities for hierarchical summarization.

## Local Search (Specific Queries)

- Query: "What did Einstein say about quantum mechanics?"
- Strategy: Entity lookup → traverse relations → retrieve text
- Uses: Entity embeddings + graph traversal

## Global Search (Broad Queries)

- Query: "What are the main research themes in this corpus?"
- Strategy: Retrieve community summaries at appropriate level
- Uses: Pre-computed hierarchical summaries

## Hybrid Search

- Combine vector similarity with graph structure
- Route based on query classification (local vs global)

---

Different query types benefit from different retrieval strategies.

## Multi-hop Query Types

- **Chain:**  $A \rightarrow B \rightarrow C$  ("Who founded the company that acquired X?")
- **Fan-out:**  $A \rightarrow \{B, C, D\}$  ("All products from company X?")
- **Fan-in:**  $\{A, B, C\} \rightarrow D$  ("What connects these entities?")

## Graph Traversal Strategies


- BFS from query entities (breadth-first search)
- Embedding-guided traversal (prioritize relevant paths)
- LLM-guided exploration (let model decide next hop)

## Challenges and Best Practices

- Path explosion with hop count – limit to 2-3 hops
- Noise accumulation – prune low-confidence edges
- Combine with vector retrieval for verification

---

Multi-hop reasoning trades coverage for precision – tune hop depth per query type.



Intelligent routing selects the optimal retrieval path.

## Architecture

- Vector store for semantic similarity
- Graph store for structural queries
- Fusion layer to combine results

## Fusion Strategies

- **Early fusion:** Combine before ranking
- **Late fusion:** Rank separately, then merge
- **Cascaded:** Vector retrieval, then graph expansion

## Implementation Options

- Separate stores: Neo4j + Pinecone
- Integrated: Neo4j with vector indexes
- LlamalIndex: PropertyGraph with vector support

---

Hybrid retrieval captures both semantic and structural relevance.

# Implementation Considerations

## Graph Storage Options

- Neo4j (native graph DB), NetworkX (in-memory), Neptune (AWS)
- Hybrid: Vector store + graph DB for combined queries

## Indexing Cost vs Query Benefit

- GraphRAG requires significant upfront processing
- Entity extraction: ~\$1-5 per 1M tokens (LLM costs)
- Summarization: Additional passes over extracted entities
- Best ROI: Large, static corpora with complex query patterns

## When to Use GraphRAG

- Multi-hop reasoning required
- Global/thematic queries common
- Corpus is relatively stable (infrequent updates)

---

GraphRAG trades indexing cost for query capability – choose based on use case.

# Required Readings

## This Week

- Edge et al. (2024). "From Local to Global: A GraphRAG Approach to Query-Focused Summarization." Microsoft Research. arXiv:2404.16130
- Pan et al. (2024). "Unifying Large Language Models and Knowledge Graphs: A Roadmap." arXiv:2306.08302

## Supplementary

- Besta et al. (2024). "Graph of Thoughts: Solving Elaborate Problems with LLMs." arXiv:2308.09687
- Gutierrez et al. (2024). "HippoRAG: Neurobiologically Inspired Long-Term Memory." arXiv:2405.14831

---

Focus on Microsoft GraphRAG paper for implementation details.

# Summary and Key Takeaways

## Key Concepts

- **Knowledge Graph:** Entity-relationship structure enabling multi-hop reasoning
- **GraphRAG:** Combine KG with vector retrieval for comprehensive search
- **Communities:** Hierarchical clustering for global query support
- **Query Routing:** Match query type to optimal retrieval strategy

## Design Principles

- Use structure when relationships matter
- Pre-compute summaries for global queries
- Route queries intelligently based on type

## Next Week

- Hallucination Prevention and Verification

---

GraphRAG = Structure + Vectors for comprehensive retrieval.