

Future Research Directions

Near-term

Better benchmarks

Production tools

Human-agent collab



Mid-term

World models

Persistent learning

Multi-agent systems



Long-term

Embodied agents

General autonomy

Agent societies

Cross-cutting: Safety | Alignment | Interpretability | Evaluation

Key insight: Progress requires interdisciplinary collaboration